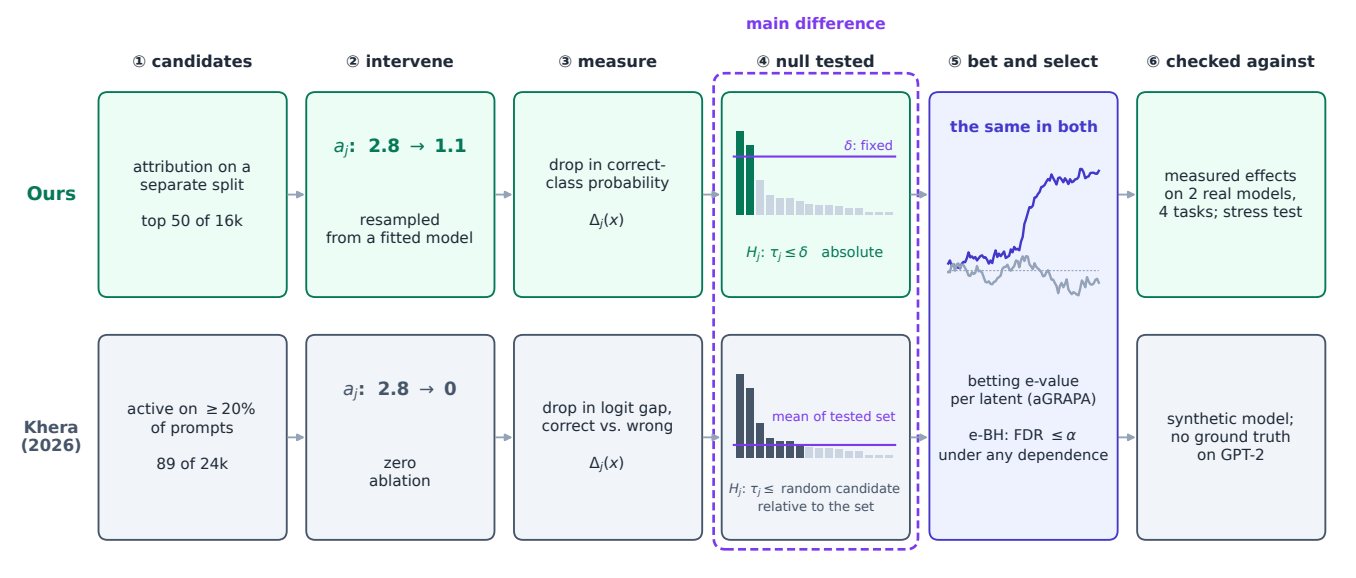
main difference
① candidates
② intervene
③ measure
④ null tested
⑤ bet and select
⑥ checked against
Ours
attribution on a separate split
top 50 of 16k
a_j: 2.8 → 1.1
resampled from a fitted model
drop in correct-class probability
$\Delta_j(x)$
δ: fixed
H_j: $\tau_j \le \delta$ absolute
the same in both
measured effects on 2 real models, 4 tasks; stress test
Khera (2026)
active on ≥ 20% of prompts
89 of 24k
a_j: 2.8 → 0
zero ablation
drop in logit gap, correct vs. wrong
$\Delta_j(x)$
mean of tested set
H_j: $\tau_j \le$ random candidate relative to the set
betting e-value per latent (aGRAPA)
e-BH: FDR $\le \alpha$ under any dependence
synthetic model; no ground truth on GPT-2