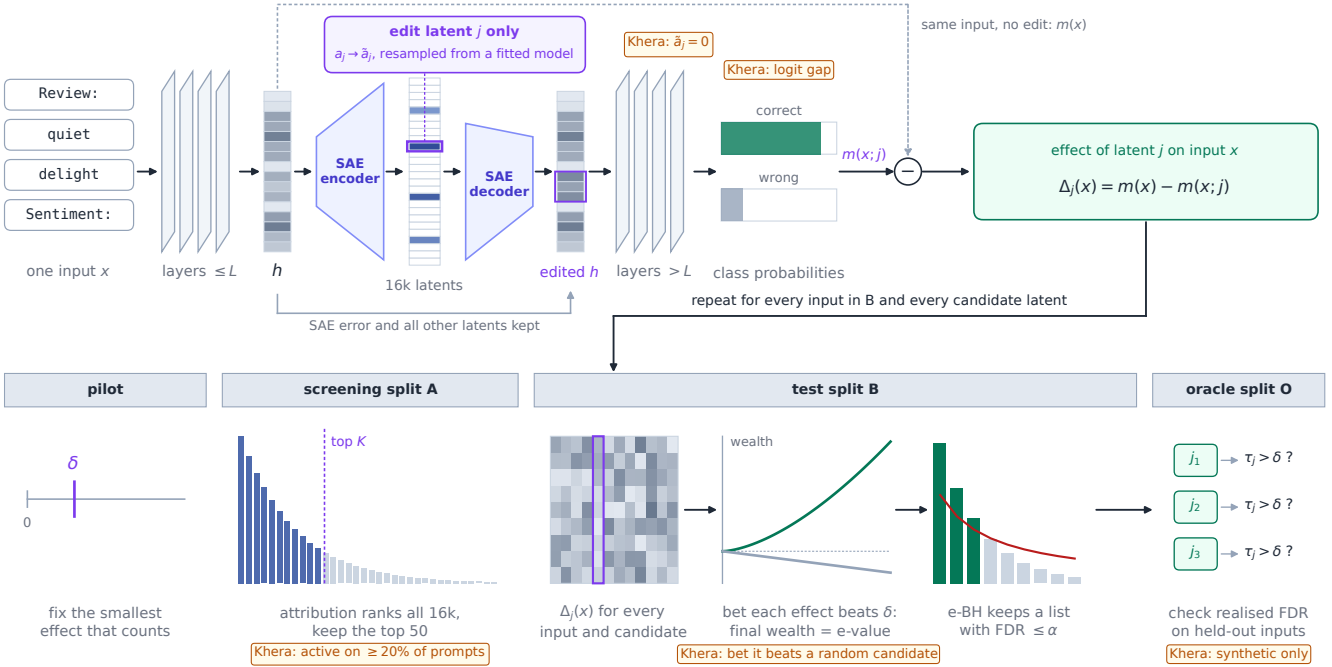
edit latent j only
$a_j \to \tilde{a}_j$, resampled from a fitted model
same input, no edit: $m(x)$
Khera: $\tilde{a}_j = 0$
Khera: logit gap
Review:
quiet
delight
Sentiment:
SAE encoder
SAE decoder
correct
wrong
$m(x; j)$
effect of latent j on input x
$\Delta_j(x) = m(x) - m(x; j)$
one input x
layers $\leq L$
h
16k latents
edited h
layers $> L$
class probabilities
repeat for every input in B and every candidate latent
SAE error and all other latents kept
pilot
screening split A
test split B
oracle split O
δ
0
top K
wealth
j_1 $\tau_j > \delta$?
j_2 $\tau_j > \delta$?
j_3 $\tau_j > \delta$?
fix the smallest effect that counts
attribution ranks all 16k, keep the top 50
Khera: active on $\geq 20\%$ of prompts
$\Delta_j(x)$ for every input and candidate
bet each effect beats δ: final wealth = e-value
Khera: bet it beats a random candidate
e-BH keeps a list with FDR $\leq \alpha$
check realised FDR on held-out inputs
Khera: synthetic only